

## ✓ Homework 6 - Convolutional Neural Networks

In this assignment, you will learn how to implement a convolutional neural network (CNN) from scratch using numpy. You will derive the derivatives yourself, implement forward and backward propagation for convolutional layers, pooling layers, and fully connected layers, and visualize the learned filters. (However, instead of using a fixed learning rate, we will call TensorFlow's implementation of Adam optimization to speed up training)

### Task 1: Load and Visualize the MNIST Data

First, let's load a subset of the MNIST dataset and visualize it. We will use this dataset to train our CNN.

```
import numpy as np
import matplotlib.pyplot as plt
from keras.datasets import mnist
from keras.utils import to_categorical
import tensorflow as tf

# Load MNIST data
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Reshape data to include channel dimension
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1)

# Normalize data
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Convert labels to one-hot encoding
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)


# Use a smaller subset of the data for faster training
x_train, y_train = x_train[:100], y_train[:100]
x_test, y_test = x_test[:200], y_test[:200]

# Visualize some samples from the dataset
def visualize_mnist_samples(X, y):
    fig, axes = plt.subplots(1, 10, figsize=(15, 15))
    for i, ax in enumerate(axes):
        ax.imshow(X[i].reshape(28, 28), cmap='gray')
        ax.set_title(np.argmax(y[i]))
```

```
ax.axis('off')
plt.show()
```

```
visualize_mnist_samples(x_train, y_train)
```

➡ Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.11490434/11490434> — 0s 0us/step



## ✓ Task 2: Implement Convolutional Layers

### Step 1: Implement the Convolution Operation

Implement the function to perform the convolution operation.

```
def conv_forward(X, W, b, stride=1, padding=1):
    (n_H_prev, n_W_prev, n_C_prev) = X.shape
    (f, f, n_C_prev, n_C) = W.shape
    n_H = int((n_H_prev - f + 2 * padding) / stride) + 1
    n_W = int((n_W_prev - f + 2 * padding) / stride) + 1
    Z = np.zeros((n_H, n_W, n_C))

    X_pad = np.pad(X, ((padding, padding), (padding, padding), (0, 0)), 'constant')

    for h in range(n_H):
        for w in range(n_W):
            for c in range(n_C):
                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f
                X_slice = X_pad[vert_start:vert_end, horiz_start:horiz_end, :]
                Z[h, w, c] = np.sum(X_slice * W[:, :, :, c]) + b[:, :, :, c]

    return Z
```

## ✓ Step 2: Implement the ReLU Activation Function

Implement the ReLU activation function.

```
def relu(Z):  
    return np.maximum(0, Z)
```

### ✓ Step 3: Implement the Max Pooling Layer (1 pt)

Implement the function to perform max pooling.

```
def pool_forward(A_prev, f=2, stride=2):  
    (n_H_prev, n_W_prev, n_C_prev) = A_prev.shape  
    n_H = int((n_H_prev - f) / stride) + 1  
    n_W = int((n_W_prev - f) / stride) + 1  
    A = np.zeros((n_H, n_W, n_C_prev))  
  
    for h in range(n_H):  
        for w in range(n_W):  
            for c in range(n_C_prev):  
                vert_start = h * stride  
                vert_end = vert_start + f  
                horiz_start = w * stride  
                horiz_end = horiz_start + f  
                A[h, w, c] = np.max(A_prev[vert_start:vert_end, horiz_start:horiz_end, c])  
  
    return A
```

### ✓ Task 3: Implement Fully Connected Layers

#### Step 1: Implement the Flatten Layer

Implement the function to flatten the input.

```
def flatten(X):  
    return X.flatten().reshape(1, -1)
```

### ✓ Step 2: Implement the Fully Connected Layer

Implement the function for the fully connected layer.

```
def fc_forward(X, W, b):  
    return np.dot(X, W) + b
```

### ✓ Step 3: Implement the Softmax Activation Function

Implement the softmax function.

```
def softmax(Z):  
    exp_Z = np.exp(Z - np.max(Z))  
    return exp_Z / np.sum(exp_Z)
```

## ✓ Task 4: Implement the Forward Propagation

Implement the forward propagation step.

```
def forward_propagation(X, params):  
    (W1, b1, W2, b2, W3, b3) = params  
  
    #Describe the computation structure of the CNN model  
    Z1 = conv_forward(X, W1, b1)  
    A1 = relu(Z1)  
    P1 = pool_forward(A1)  
    F1 = flatten(P1)  
    Z2 = fc_forward(F1, W2, b2)  
    A2 = relu(Z2)  
    Z3 = fc_forward(A2, W3, b3)  
    A3 = softmax(Z3)  
  
    return Z1, A1, P1, F1, Z2, A2, Z3, A3
```

## ✓ Task 5: Compute the Loss

Implement the function to compute the cross-entropy loss.

```
def compute_loss(y, y_pred):  
    return -np.log(y_pred[0, y])
```

## ✓ Task 6: Implement the Backward Propagation

Compute the gradients using backward propagation.

```
def backward_propagation(X, y, params, cache):  
    (W1, b1, W2, b2, W3, b3) = params  
    (Z1, A1, P1, F1, Z2, A2, Z3, A3) = cache  
  
    dZ3 = A3  
    dZ3[0, y] -= 1
```

```

dW3 = np.dot(A2.T, dZ3)
db3 = np.sum(dZ3, axis=0, keepdims=True)
dA2 = np.dot(dZ3, W3.T)

dZ2 = dA2 * (Z2 > 0)
dW2 = np.dot(F1.T, dZ2)
db2 = np.sum(dZ2, axis=0, keepdims=True)
dF1 = np.dot(dZ2, W2.T)

dP1 = dF1.reshape(P1.shape)
dA1 = np.zeros_like(A1)
for h in range(P1.shape[0]):
    for w in range(P1.shape[1]):
        for c in range(P1.shape[2]):
            vert_start = h * 2
            vert_end = vert_start + 2
            horiz_start = w * 2
            horiz_end = horiz_start + 2
            mask = (A1[vert_start:vert_end, horiz_start:horiz_end, c] == np.max(A1[vert_start:vert_end, horiz_start:horiz_end, c]))
            dA1[vert_start:vert_end, horiz_start:horiz_end, c] += mask * dP1[h, w, c]

dZ1 = dA1 * (Z1 > 0)
dW1 = np.zeros_like(W1)
db1 = np.zeros_like(b1)
X_pad = np.pad(X, ((1, 1), (1, 1), (0, 0)), 'constant')
for h in range(dZ1.shape[0]):
    for w in range(dZ1.shape[1]):
        for c in range(dZ1.shape[2]):
            vert_start = h * 1
            vert_end = vert_start + W1.shape[0]
            horiz_start = w * 1
            horiz_end = horiz_start + W1.shape[1]
            dW1[:, :, :, c] += dZ1[h, w, c] * X_pad[vert_start:vert_end, horiz_start:horiz_end, :]
            db1[:, :, :, c] += dZ1[h, w, c]

return dW1, db1, dW2, db2, dW3, db3

```

## ✓ Task 7: Update Parameters Using Gradient Descent

Update the parameters using gradient descent.

```

def update_parameters(params, grads, learning_rate):
    (W1, b1, W2, b2, W3, b3) = params
    (dW1, db1, dW2, db2, dW3, db3) = grads

    # Convert NumPy arrays to TensorFlow variables
    W1 = tf.Variable(W1)
    b1 = tf.Variable(b1)

```

```

W2 = tf.Variable(W2)
b2 = tf.Variable(b2)
W3 = tf.Variable(W3)
b3 = tf.Variable(b3)

optimizer.apply_gradients(zip([dW1, db1, dW2, db2, dW3, db3], [W1, b1, W2, b2, W3, b3]))

# Convert TensorFlow variables back to NumPy arrays for consistency
W1 = W1.numpy()
b1 = b1.numpy()
W2 = W2.numpy()
b2 = b2.numpy()
W3 = W3.numpy()
b3 = b3.numpy()

#W1 -= learning_rate * dW1
#b1 -= learning_rate * db1
#W2 -= learning_rate * dW2
#b2 -= learning_rate * db2
#W3 -= learning_rate * dW3
#b3 -= learning_rate * db3

return W1, b1, W2, b2, W3, b3

```

## ✓ Task 8: Train the Model

Train the model using gradient descent (with Adam optimization) and plot the loss over iterations.

```

def initialize_parameters_conv(filter_shape):
    W = np.random.randn(*filter_shape) * 0.01
    b = np.zeros((1, 1, 1, filter_shape[3]))
    return W, b

def initialize_parameters_fc(shape):
    W = np.random.randn(shape[0], shape[1]) * 0.01
    b = np.zeros((1, shape[1]))
    return W, b

input_dim = (28, 28, 1)
conv_dim = (3, 3, 1, 8)
fc1_dim = (8 * 14 * 14, 100)
fc2_dim = (100, 10)
W1, b1 = initialize_parameters_conv(conv_dim)
W2, b2 = initialize_parameters_fc(fc1_dim)
W3, b3 = initialize_parameters_fc(fc2_dim)

# Train the model
num_iterations = 50

```

```

learning_rate = 0.01
optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
loss_history = []

for i in range(num_iterations):
    loss = 0

    grad_dW1 = np.zeros_like(W1)
    grad_db1 = np.zeros_like(b1)
    grad_dW2 = np.zeros_like(W2)
    grad_db2 = np.zeros_like(b2)
    grad_dW3 = np.zeros_like(W3)
    grad_db3 = np.zeros_like(b3)

    params = (W1, b1, W2, b2, W3, b3)

    for j in range(len(x_train)):
        X = x_train[j]
        y = np.argmax(y_train[j])

        # Forward propagation
        Z1, A1, P1, F1, Z2, A2, Z3, A3 = forward_propagation(X, params)
        cache = (Z1, A1, P1, F1, Z2, A2, Z3, A3)

        # Compute loss
        loss += compute_loss(y, A3)

        # Backward propagation
        dW1, db1, dW2, db2, dW3, db3 = backward_propagation(X, y, params, cache)

        # Accumulate gradients across training examples
        grad_dW1 += dW1
        grad_db1 += db1
        grad_dW2 += dW2
        grad_db2 += db2
        grad_dW3 += dW3
        grad_db3 += db3

    # Reinitialize the optimizer for the new set of variables
    optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)

    # Update parameters
    grads = (grad_dW1, grad_db1, grad_dW2, grad_db2, grad_dW3, grad_db3)
    W1, b1, W2, b2, W3, b3 = update_parameters(params, grads, learning_rate)
    params = (W1, b1, W2, b2, W3, b3)
    print(f'Iteration {i}: Loss = {loss}')
    loss_history.append(loss)

# Plot the loss over iterations
plt.plot(loss_history)
plt.xlabel('Iteration')

```


```
plt.ylabel('Loss')  
plt.title('Loss over Iterations')  
plt.show()
```



↳ <ipython-input-2-89709a2af507>:18: DeprecationWarning: Conversion of an array with ndim  
Z[h, w, c] = np.sum(X\_slice \* W[:, :, :, c]) + b[:, :, :, c]  
Iteration 0: Loss = 230.2617107752144  
Iteration 1: Loss = 229.5248364596287  
Iteration 2: Loss = 222.95291962999062  
Iteration 3: Loss = 207.208246067131  
Iteration 4: Loss = 182.37067347743118  
Iteration 5: Loss = 148.9461666135724  
Iteration 6: Loss = 145.99463211437805  
Iteration 7: Loss = 135.27252052242235  
Iteration 8: Loss = 119.53513185562568  
Iteration 9: Loss = 111.44066421297269  
Iteration 10: Loss = 98.98291075581086  
Iteration 11: Loss = 99.21221376660549  
Iteration 12: Loss = 87.66154187482398  
Iteration 13: Loss = 88.59188937328139  
Iteration 14: Loss = 73.06703366433213  
Iteration 15: Loss = 81.8091446712278  
Iteration 16: Loss = 66.89382075490425  
Iteration 17: Loss = 83.09952330792639  
Iteration 18: Loss = 60.25588725363192  
Iteration 19: Loss = 82.79499092872382  
Iteration 20: Loss = 52.35144130583416  
Iteration 21: Loss = 74.77150662027081  
Iteration 22: Loss = 49.512990017628624  
Iteration 23: Loss = 68.21102359017489  
Iteration 24: Loss = 41.57148912476719  
Iteration 25: Loss = 57.5919217009696  
Iteration 26: Loss = 35.261159311704866  
Iteration 27: Loss = 49.469785310501834  
Iteration 28: Loss = 30.94068266395037  
Iteration 29: Loss = 45.74459624023772  
Iteration 30: Loss = 29.077113700044205  
Iteration 31: Loss = 45.559050778240255  
Iteration 32: Loss = 26.887049753851283  
Iteration 33: Loss = 38.725517206751725  
Iteration 34: Loss = 23.83562798071274  
Iteration 35: Loss = 31.44796266690198  
Iteration 36: Loss = 20.252678594449385  
Iteration 37: Loss = 34.2159944067301  
Iteration 38: Loss = 20.04980042787262  
Iteration 39: Loss = 26.959537637175373  
Iteration 40: Loss = 18.067003005901118  
Iteration 41: Loss = 24.880366600354368  
Iteration 42: Loss = 17.31248898881672  
Iteration 43: Loss = 20.305227883025115  
Iteration 44: Loss = 15.472796024947641  
Iteration 45: Loss = 17.692684442898205  
Iteration 46: Loss = 13.660007063378066  
Iteration 47: Loss = 14.414953729045305  
Iteration 48: Loss = 12.489689893185469  
Iteration 49: Loss = 11.699063992492695

Loss over Iterations





## ✓ Task 9: Visualize the Learned Filters

Visualize the filters learned by the first convolutional layer.

```
def visualize_filters(W1):  
    num_filters = W1.shape[-1]  
    fig, axes = plt.subplots(1, num_filters, figsize=(20, 20))  
    for i in range(num_filters):  
        axes[i].imshow(W1[:, :, :, i], cmap='gray')  
        axes[i].axis('off')  
    plt.show()  
  
visualize_filters(W1)
```