
```

import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping
import matplotlib.pyplot as plt
import numpy as np

# Normalize function: Subtract mean and divide by standard deviation
def normalize(x):
    mean = np.mean(x, axis=(0, 1, 2), keepdims=True) # Mean across height, width, and color channels
    std = np.std(x, axis=(0, 1, 2), keepdims=True) # Standard deviation across height, width, and color channels
    return (x - mean) / (std + 1e-7) # Adding small epsilon to avoid division by zero

# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Normalize data to [0, 1] range
x_train = normalize(x_train)
x_test = normalize(x_test)

# One-hot encode the labels
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

# Lowers learning rate if not improving for 3 epochs
lr_schedule = tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e-6)

# Stops early if not improving for 5 epochs
early_stopping = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

# Residual network architecture
def residual_block(x, filters):
    skip = x

    # If the number of filters doesn't match, apply a 1x1 convolution
    if x.shape[-1] != filters:
        skip = layers.Conv2D(filters, (1, 1), strides=1, padding='same')(x)

    x = layers.Conv2D(filters, (3, 3), padding='same', activation='relu')(x)
    x = layers.Conv2D(filters, (3, 3), padding='same', activation='relu')(x)

    # Add input to convolutional outputs (makes it residual)
    x = layers.add([x, skip])

    x = layers.BatchNormalization()(x)
    x = layers.ReLU()(x)
    return x

# Adam optimizer
optimizer = tf.keras.optimizers.Adam(learning_rate=1e-4)

# Input layer
inputs = layers.Input(shape=(32, 32, 3))

# Initial Conv Layer
model_pipe = layers.Conv2D(32, (3, 3), padding='same', activation='relu')(inputs)
model_pipe = layers.BatchNormalization()(model_pipe)
model_pipe = layers.MaxPooling2D((2, 2))(model_pipe)

```

```

model_pipe = layers.MaxPooling2D((2, 2))(model_pipe)

# Residual Layers, double filters each time
model_pipe = residual_block(model_pipe, 64)
model_pipe = layers.MaxPooling2D((2, 2))(model_pipe)

model_pipe = residual_block(model_pipe, 128)
model_pipe = layers.MaxPooling2D((2, 2))(model_pipe)

model_pipe = residual_block(model_pipe, 256)
model_pipe = layers.MaxPooling2D((2, 2))(model_pipe)

# Flatten output from Residual Layers for Fully Connected Layers
model_pipe = layers.Flatten()(model_pipe)

# Fully Connected Layers, use dropout to prevent overfitting
model_pipe = layers.Dense(256, activation='relu')(model_pipe)
model_pipe = layers.Dropout(0.5)(model_pipe)
model_pipe = layers.Dense(128, activation='relu')(model_pipe)
model_pipe = layers.Dropout(0.5)(model_pipe)

# Use softmax to create probabilities for each class
model_pipe = layers.Dense(10, activation='softmax')(model_pipe)

# Create model based on pipeline and inputs
model = models.Model(inputs, model_pipe)

# Compile model using categorical cross entropy due to multi-class classification task
model.compile(optimizer=optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()

```



Layer (type)	Output Shape	Param #	Connected to
input_layer (InputLayer)	(None, 32, 32, 3)	0	-
conv2d (Conv2D)	(None, 32, 32, 32)	896	input_layer[0][0]
batch_normalization (BatchNormalization)	(None, 32, 32, 32)	128	conv2d[0][0]
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0	batch_normalization[0...
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18,496	max_pooling2d[0][0]
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36,928	conv2d_2[0][0]
conv2d_1 (Conv2D)	(None, 16, 16, 64)	2,112	max_pooling2d[0][0]
add (Add)	(None, 16, 16, 64)	0	conv2d_3[0][0], conv2d_1[0][0]
batch_normalization_1 (BatchNormalization)	(None, 16, 16, 64)	256	add[0][0]
re_lu (ReLU)	(None, 16, 16, 64)	0	batch_normalization_1...
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0	re_lu[0][0]
conv2d_5 (Conv2D)	(None, 8, 8, 128)	73,856	max_pooling2d_1[0][0]
conv2d_6 (Conv2D)	(None, 8, 8, 128)	147,584	conv2d_5[0][0]
conv2d_4 (Conv2D)	(None, 8, 8, 128)	8,320	max_pooling2d_1[0][0]
add_1 (Add)	(None, 8, 8, 128)	0	conv2d_6[0][0], conv2d_4[0][0]
batch_normalization_2 (BatchNormalization)	(None, 8, 8, 128)	512	add_1[0][0]
re_lu_1 (ReLU)	(None, 8, 8, 128)	0	batch_normalization_2...
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0	re_lu_1[0][0]
conv2d_8 (Conv2D)	(None, 4, 4, 256)	295,168	max_pooling2d_2[0][0]
conv2d_9 (Conv2D)	(None, 4, 4, 256)	590,080	conv2d_8[0][0]
conv2d_7 (Conv2D)	(None, 4, 4, 256)	33,024	max_pooling2d_2[0][0]
add_2 (Add)	(None, 4, 4, 256)	0	conv2d_9[0][0], conv2d_7[0][0]
batch_normalization_3 (BatchNormalization)	(None, 4, 4, 256)	1,024	add_2[0][0]
re_lu_2 (ReLU)	(None, 4, 4, 256)	0	batch_normalization_3...
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 256)	0	re_lu_2[0][0]
flatten (Flatten)	(None, 1024)	0	max_pooling2d_3[0][0]
dense (Dense)	(None, 256)	262,400	flatten[0][0]
dropout (Dropout)	(None, 256)	0	dense[0][0]

dense_1 (Dense)	(None, 128)	32,896	dropout[0][0]
dropout_1 (Dropout)	(None, 128)	0	dense_1[0][0]
dense_2 (Dense)	(None, 10)	1,290	dropout_1[0][0]

Total params: 1,504,970 (5.74 MB)

Trainable params: 1,504,010 (5.74 MB)

Train the model

```
history = model.fit(x_train, y_train, epochs=50, batch_size=64,
                    validation_data=(x_test, y_test),
                    callbacks=[early_stopping, lr_schedule])
```

Evaluate the model

```
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print(f"Test accuracy: {test_acc:.4f}")
```

Plot training history

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```