

✓ MNIST Dimensionality Reduction [60 Points]

First, we are going to explore two different strategies to reduce the dimensionality of digit images. One, we shall build a CNN-based autoencoder. After, we shall reduce the dimensionality of images using PCA. Again, **avoid any unnecessary loops**, using [broadcasting](#), and writing [vectorized](#) code.

✓ Importing Libraries and Data

You may have to select GPU acceleration in Edit > Notebook Settings. No worries if you do not have GPU available, it may just take longer.

```
import tensorflow as tf
device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    print('GPU device not found')
else:
    print('Found GPU at: {}'.format(device_name))
```

```
Found GPU at: /device:GPU:0
Found GPU at: /device:GPU:0
```

```
from tensorflow import keras
import numpy as np
import matplotlib.pyplot as plt
plt.style.use('seaborn-v0_8-white')
```

```
import tensorflow_datasets as tfds
(ds_train_org, ds_test_org), ds_info = tfds.load(
    'mnist',
    split=['train', 'test'],
    shuffle_files=True,
    as_supervised=True,
    with_info=True,
)
```

```
Downloading and preparing dataset 11.06 MiB (download: 11.06 MiB, generated: 21.00 MiB, total: 32.06 MiB) to /root/tensorflow_datasets/mnist
DI Completed...: 100% 5/5 [00:05<00:00, 1.28s/ file]
Dataset mnist downloaded and prepared to /root/tensorflow_datasets/mnist/3.0.1. Subsequent calls will reuse this data.
```

```
def normalize_img(image, label):
    """Normalizes images: `uint8` -> `float32`."""
    return tf.cast(image, tf.float32) / 255.0

def normalize_img_label(image, label):
    """Normalizes images: `uint8` -> `float32`."""
    return tf.cast(image, tf.float32) / 255.0, label

ds_train = ds_train_org.map(
    normalize_img, num_parallel_calls=tf.data.AUTOTUNE)
ds_test_labels = ds_test_org.map(
    normalize_img_label, num_parallel_calls=tf.data.AUTOTUNE)
```

✓ MNIST CNN Autoencoder [35 Points]

First, we are going to build a neural network autoencoder.

[10 points] Let's define the autoencoder neural network.

```
N_batch = 200 # batch size
code_size = 32 # size of autoencoder codes
hid_size = 128 # number of hidden units in network layers
```

```
# encode_nnet should be keras model that has 3 Conv2D layers with hid_size channels
# the filter sizes should be 4 by 4, with a stride of 2 in each dimension
# use a relu activation in the first two layers, none in the last
```

```
encode_nnet = keras.models.Sequential([
    keras.layers.Conv2D(hid_size, 4, strides=(2, 2), activation='relu'),
    keras.layers.Conv2D(hid_size, 4, strides=(2, 2), activation='relu'),
    keras.layers.Conv2D(code_size, 4, strides=(2, 2))])
```

```
encode_nnet.build([N_batch, 28, 28, 1])
```

```
# encode_nnet.output.shape # should be (N_batch, 1, 1, code_size)!
```

Next we are going to use a very similar type of convolutional operation, the [transposed convolution](#), to upsample and reconstruct the original image starting with the encoder codes.

```
# decode_nnet should be keras model that has 3 Conv2DTranspose layers with
# hid_size channels, the filter sizes should be 5 by 5 for the first two layers
# and 4 by 4 for the last, with a stride of 2 in each dimension.
# use a relu activation in the first two layers, sigmoid in the last.
```

```
decode_nnet = keras.models.Sequential([
    keras.layers.Conv2DTranspose(hid_size, 5, strides=(2, 2), activation='relu'),
    keras.layers.Conv2DTranspose(hid_size, 5, strides=(2, 2), activation='relu'),
    keras.layers.Conv2DTranspose(1, 4, strides=(2, 2), activation='sigmoid')
]) # TODO
```

```
# decode_nnet.build([200,1,1,32])
```

```
# decode_nnet.output.shape # should be (N_batch, 28, 28, 1)
```

[15 points] Let's implement the loss for batches.

```
def loss(X, encoder, decoder):
    """
    Args:
        X: tensor of input images (N_batch, 28, 28, 1)
        encoder: function that takes in a batch of images and outputs codes
        decoder: function that takes in a batch of codes and outputs reconstructions
    Returns:
        loss: mean of the sum of squared errors (Euclidean distance squared) for
            reconstructions
    """
    # recons_loss should be a mean of the --sum of squared errors-- for
    # reconstructions, averaged over all instances in the batch
    # MAKE SURE TO USE TENSORFLOW FUNCTIONS AND OPERATIONS (IE NOT NUMPY)
    # Hint: make sure to use reduce_mean and reduce_sum over the correct axes.

    encoded_X = encoder(X)
    reconstructed_X = decoder(encoded_X)

    squared_errors = tf.square(X - reconstructed_X)
    sum_squared_errors = tf.reduce_sum(squared_errors, axis=[1, 2, 3])

    recons_loss = tf.reduce_mean(sum_squared_errors) # TODO

    return recons_loss

def grad(X, enc_nnet, dec_nnet):
    with tf.GradientTape() as tape:
        loss_value = loss(X, enc_nnet, dec_nnet)

    return tape.gradient(loss_value, enc_nnet.weights+dec_nnet.weights)
```

```
optimizer = tf.keras.optimizers.Adam(learning_rate=0.00001)
# nepochs = number of times to run through the data.
nepochs = 25 # Could probably do better with more epochs, but should suffice
for i in range(nepochs):
    for B_X in ds_train.batch(N_batch, drop_remainder=True):
        grads = grad(B_X, encode_nnet, decode_nnet)
```

```
optimizer.apply_gradients(zip(grads, encode_nnet.weights+decode_nnet.weights)) # SGD-type update (w/ Adam)
print("Loss at epoch {:03d}: {:.3f}".format(i, loss(B_X, encode_nnet, decode_nnet)))
```

```
↳ Loss at epoch 000: 98.774
Loss at epoch 001: 63.028
Loss at epoch 002: 54.175
Loss at epoch 003: 53.287
Loss at epoch 004: 51.995
Loss at epoch 005: 48.933
Loss at epoch 006: 44.586
Loss at epoch 007: 39.727
Loss at epoch 008: 36.528
Loss at epoch 009: 34.390
Loss at epoch 010: 32.494
Loss at epoch 011: 30.693
Loss at epoch 012: 29.202
Loss at epoch 013: 28.004
Loss at epoch 014: 27.021
Loss at epoch 015: 26.169
Loss at epoch 016: 25.416
Loss at epoch 017: 24.717
Loss at epoch 018: 24.022
Loss at epoch 019: 23.330
Loss at epoch 020: 22.647
Loss at epoch 021: 22.001
Loss at epoch 022: 21.389
Loss at epoch 023: 20.824
Loss at epoch 024: 20.324
```

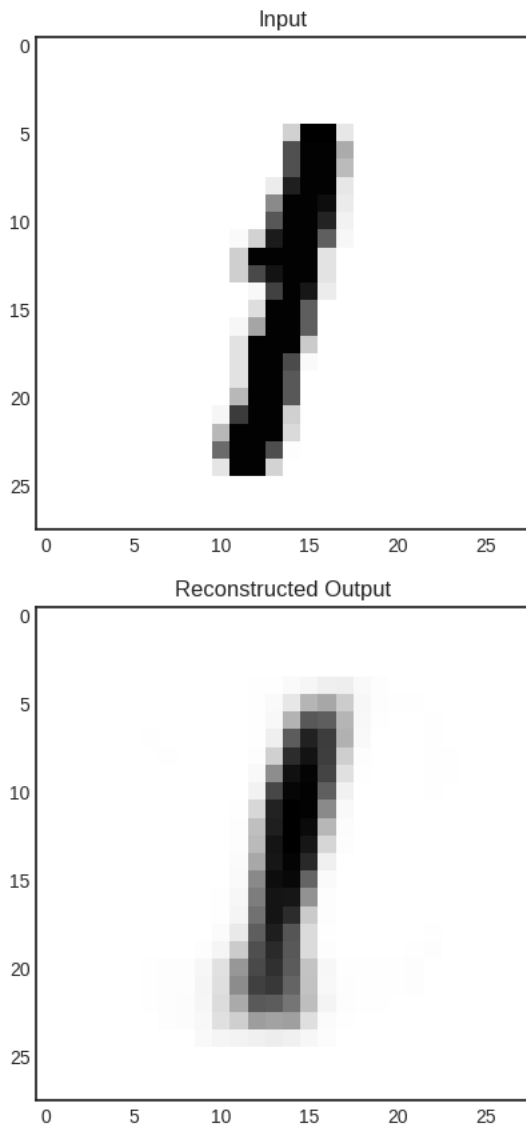
[5 points] Let's visualize the output for an instance in the batch.

```
B_X_plot = B_X
nnet_output = decode_nnet(encode_nnet(B_X)) # TODO: get N_batch x 28 x 28 x 1 tensor of reconstructions
rind = np.random.randint(0, N_batch) # Select a random instance in batch

plt.figure()
plt.title('Input')
plt.imshow(np.reshape(B_X_plot[rind], (28, 28)))

plt.figure()
plt.title('Reconstructed Output')
x_recon = nnet_output[rind]
plt.imshow(np.reshape(x_recon, (28, 28)))
```


 <matplotlib.image.AxesImage at 0x7b6ce0bdf10>



[5 points] Let's compute the accuracy over the test set.

```
test_batch_MSEs = []
for B_X, B_Ylbls in ds_test_labels.batch(N_batch, drop_remainder=True):
    nnet_output = decode_nnet(encode_nnet(B_X))
    squared_errors = tf.square(B_X - nnet_output)
    sum_squared_errors = tf.reduce_sum(squared_errors, axis=[1, 2, 3])
    test_batch_MSEs.append(tf.reduce_mean(sum_squared_errors)) # TODO: append the MSE for the batch

print('Mean Sum of Squared Errors on test set: {}'.format(np.mean(test_batch_MSEs)))
```

 Mean Sum of Squared Errors on test set: 20.757062911987305

✓ PCA [25 Points]

Next, we are going to learn a linear autoencoder with PCA, utilizing scikit-learn's implementation.

[10 points] Let's fit a PCA model on the MNIST data.

```
from sklearn.decomposition import PCA
pca = PCA(n_components=code_size)

# Hint: list(ds_train.as_numpy_iterator()) is very useful
Xdesign_mat = []
```

```
Xdesign_mat = np.stack(list(ds_train.as_numpy_iterator()), axis=0).reshape(-1,784)
pca.fit_transform(Xdesign_mat)
```

```
array([[ -2.7357507 , -0.40374228, -0.32074684, ..., -0.37409383,
        -0.3000402  ,  0.9021928  ],
       [ -3.9322762 ,  1.6692578 , -0.44134945, ..., -0.98440295,
        -0.3824799 ,  0.51441216 ],
       [  4.5570955 ,  0.9185379 , -3.938978 , ..., -0.05938989,
        -0.25504458, -0.42250964 ],
       ...,
       [  2.9310338 , -0.03648547, -0.22359508, ..., -0.26937854,
        0.83677745,  0.04880274 ],
       [ -3.1157043 ,  0.46734783,  0.80377257, ..., -0.18952505,
        0.06372073,  0.5626818  ],
       [ -0.4871719 , -0.28548703, -1.1456549 , ...,  1.292084  ,
        1.9250314 ,  0.37560377 ]], dtype=float32)
```


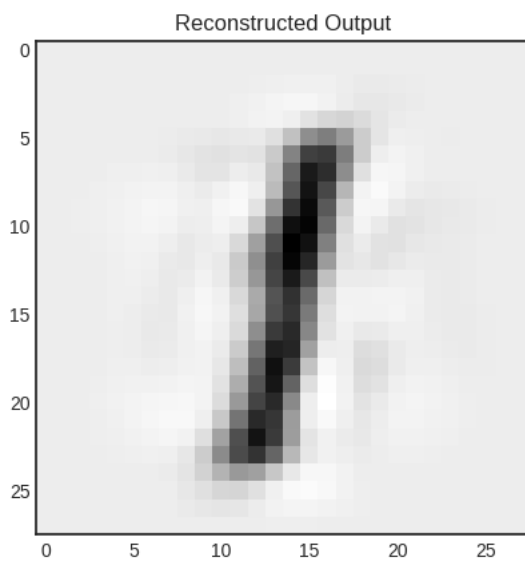

[5 points] Let's compute MINST reconstructions using PCA.

```
B_X_plot = tf.reshape(B_X_plot, [B_X_plot.shape[0], -1]).numpy()
codes = pca.transform(B_X_plot)
recons = np.reshape(pca.inverse_transform(codes), (B_X_plot.shape[0], 28, 28))
```

```
plt.figure()
plt.title('Input')
plt.imshow(np.reshape(B_X_plot[rind], (28, 28)))
```

```
plt.figure()
plt.title('Reconstructed Output')
x_recon = recons[rind]
plt.imshow(np.reshape(x_recon, (28, 28)))
```

```
plt.figure()
plt.title('Truncated Reconstructed Output')
x_recon[x_recon<0.0] = 0.0
x_recon[x_recon>1.0] = 1.0
plt.imshow(np.reshape(x_recon, (28, 28)))
```

 matplotlib.image.AxesImage at 0x700c5091c4107


[10 points] Let's compute MNIST PCA reconstructions errors on the MNIST test set.

```
test_batch_MSEs_pca = []
for B_X, B_Ybls in ds_test_labels.batch(N_batch, drop_remainder=True):
    # YOU must use loss(X, encoder, decoder) here!!
```

```
# Do *not* truncate reconstructions.
# Hint: lambda functions...
encoder = lambda x: pca.transform(tf.reshape(x, [x.shape[0], -1]).numpy())
decoder = lambda x: tf.convert_to_tensor(pca.inverse_transform(x).reshape(-1, 28, 28, 1))
test_batch_MSEs_pca.append(loss(B_X, encoder, decoder)) # TODO: append the MSE for the batch

print('Mean Sum of Squared Errors on test set: {}'.format(
    np.mean(test_batch_MSEs_pca)))
```

➞ Mean Sum of Squared Errors on test set: 13.193336486816406

Not too bad for a linear model!!

✓ Kmeans from Scratch [40 Points]


We are now going to implement and visualize kmeans. In addition to getting some more intuition and understanding about kmeans, we are going to be efficient about our implementation. This means **avoid any unnecessary loops**, using [broadcasting](#), and writing [vectorized](#) code.

First, let's generate data.

```
np.random.seed(8675309) # Feel free to fiddle with seed when playing around
X1 = np.concatenate([0.5*np.random.randn(100, 1)-2.0, 0.5*np.random.randn(100, 1)-2.0], 1)
X2 = np.concatenate([0.5*np.random.randn(100, 1)+-2.0, 0.5*np.random.randn(100, 1)], 1)
X3 = np.concatenate([0.5*np.random.randn(100, 1)+2.0, 0.1*np.random.randn(100, 1)-1.0], 1)
X4 = np.concatenate([0.5*np.random.randn(100, 1)+2.0, 0.1*np.random.randn(100, 1)+1.0], 1)
X_all = np.concatenate([X1, X2, X3, X4], 0)
mus_init = np.random.randn(4, 2)
```

```
plt.scatter(X_all[:, 0], X_all[:, 1])
```


➞ <matplotlib.collections.PathCollection at 0x7b6ce0b36ec0>



Visualizing the "ground truth" clusters.

```
plt.scatter(X1[:, 0], X1[:, 1])
plt.scatter(X2[:, 0], X2[:, 1])
plt.scatter(X3[:, 0], X3[:, 1])
plt.scatter(X4[:, 0], X4[:, 1])
plt.scatter(mus_init[:, 0], mus_init[:, 1], c='black')
```

 <matplotlib.collections.PathCollection at 0x7b6c595c77f0>



[15 points] Let's implement the "distortion" loss that kmeans is minimizing. **Avoid any loops.**

```
def kmeans_distortion(X, mus, zs_onehot):
    """
    Args:
        X: N x d matrix of the points to cluster
        mus: K x d matrix of means
        zs_onehot: N x K matrix of one hot indicators of assignments.
            I.e. zs_onehot[i, j] = 1 if point i is in clust j, 0 otherwise.
    Returns:
        distortion: scalar of loss that Kmeans is trying to minimize
    """
    # hint: use diffs_cluster[i, j] = ||x_i - mu_j||^2 matrix
    # other hint: diffs_cluster[i, j] = ||x_i||^2 - 2 x_i^mu_j + ||mu_j||^2

    X_norm_sq = np.sum(X**2, axis=1, keepdims=True)
    mus_norm_sq = np.sum(mus**2, axis=1)
    diffs_cluster = X_norm_sq - 2 * np.dot(X, mus.T) + mus_norm_sq
    distortion = np.sum(zs_onehot * diffs_cluster)

    return distortion
```

[15 points] Let's now code kmeans. It's a simple algorithm that can actually be coded in a handful of lines. Let's be efficient, and **avoid any other loops.**

```
def kmeans_cluster(X, mus, iters):
    """
    Args:
        X: N x d matrix of the points to cluster
        mus: K x d matrix of initial means
        iters: integer number of iterations to run kmeans updates
    Returns:
        mus: K x d matrix of means
        zs_onehot: N x K matrix of one hot indicators of assignments.
            I.e. zs_onehot[i, j] = 1 if point i is in clust j, 0 otherwise.
        distortions: iters length array of loss that Kmeans is trying to minimize
    """
    K = mus.shape[0]
    eyeK = np.eye(K)

    distortions = []
    for t in range(iters):
        # get distances
        # diffs_cluster[i, j] = ||x_i - mu_j||^2
        X_norm_sq = np.sum(X**2, axis=1, keepdims=True)
        mus_norm_sq = np.sum(mus**2, axis=1)
        diffs_cluster = X_norm_sq - 2 * np.dot(X, mus.T) + mus_norm_sq # should be N x k matrix

        # assignments
        # zs[i] \in {1, .., K}, cluster x_i is in
```



```

zs = np.argmin(diffs_cluster, axis=1) # N length array
# zs_onehot[i, j] = x_i \in cluster_j
zs_onehot = eyeK[zs] # should be N x k matrix

# record distortion after new assignments
distortions.append(kmeans_distortion(X, mus, zs_onehot))

# plot assignments
plt.figure()
plt.title('Iter {}'.format(t))
plt.scatter(X[:, 0], X[:, 1], c=zs)
plt.scatter(mus[:, 0], mus[:, 1], c='red')

# update mus
# mus[j, 1] = 1-th dimension of cluster j
for j in range(K):
    if np.sum(zs_onehot[:, j]) > 0: # Avoid division by zero
        mus[j] = np.sum(X * zs_onehot[:, j][:, np.newaxis], axis=0) / np.sum(zs_onehot[:, j])
# TODO: should be K x d matrix

# plot new means
plt.scatter(mus[:, 0], mus[:, 1], c='orange')

# record distortion after new means
distortions.append(kmeans_distortion(X, mus, zs_onehot))

plt.figure()
plt.title('Distortions')
plt.plot(distortions)
return mus, zs_onehot, distortions

```

[+ Code](#)
[+ Text](#)

```

plt.style.use('default')

mus_kmeans, zs_onehot_kmeans, distortions_kmeans = kmeans_cluster(X_all, mus_init, 6)

```